

Practical Machine Learning - Course Final Project

Alberto A. Caeiro Jr

February 18, 2015

Executive Summary & Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways.

The goal of this project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

About the Data

The data for this project are available here: . Training dataset : “<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>” (<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>)” . Testing dataset:

“<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv> (<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>)” The data for this project come from this source: <http://groupware.les.inf.puc-rio.br/har> (<http://groupware.les.inf.puc-rio.br/har>). If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment. So let’s have a look on the dataset and on the classe variable.

Data Processing

Obtaining and Cleaning the data

```

# Enabling Multi Core for modeling processing
library(doMC)
  registerDoMC(cores = 2)

#Loading used libraries
library(caret);library(klaR); library(rpart)
library(randomForest); library(gbm)

#setting the seed for reproducible computation
set.seed(12345)

#setting the working directory folder
setwd("~/Developer/Data Science Specialization/Practical Machine Learning/Project")

# loading both testing and training dataset (considering both files were already downloaded)
trainFile <- "./pml-training.csv"
training <- read.csv(file=trainFile, header=TRUE, sep="," , na.strings=c("NA","#DIV/0!",""))
testFile <- "./pml-testing.csv"
testing <- read.csv(file=testFile, header=TRUE, sep="," , na.strings=c("NA","#DIV/0!",""))

# Summary for the training predictors and outcome
str(training)

```

```

## 'data.frame':   19622 obs. of  160 variables:
## $ X : int  1 2 3 4 5 6 7 8 9 10 ...
## $ user_name : Factor w/ 6 levels "adelmo","carlitos",...: 2 2 2 2 2 2 2 2 2 2 ...
## $ raw_timestamp_part_1 : int  1323084231 1323084231 1323084231 1323084232 1323084232 1323084232 1323084232 1323084232 1323084232 1323084232 ...
## $ raw_timestamp_part_2 : int  788290 808298 820366 120339 196328 304277 368296 440390 484323 484434 ...
## $ cvtd_timestamp : Factor w/ 20 levels "02/12/2011 13:32",...: 9 9 9 9 9 9 9 9 9 9 ...
## $ new_window : Factor w/ 2 levels "no","yes": 1 1 1 1 1 1 1 1 1 1 ...
## $ num_window : int  11 11 11 12 12 12 12 12 12 12 ...
## $ roll_belt : num  1.41 1.41 1.42 1.48 1.48 1.45 1.42 1.42 1.43 1.45 ...
## $ pitch_belt : num  8.07 8.07 8.07 8.05 8.07 8.06 8.09 8.13 8.16 8.17 ...

```

```

## $ yaw_belt          : num  -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 -94.4 ...
## $ total_accel_belt  : int    3 3 3 3 3 3 3 3 3 3 ...
## $ kurtosis_roll_belt : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_belt : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_belt  : logi   NA NA NA NA NA NA NA ...
## $ skewness_roll_belt : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_roll_belt.1 : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_belt  : logi   NA NA NA NA NA NA NA ...
## $ max_roll_belt      : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_belt     : int    NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_belt       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_belt      : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_belt     : int    NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_belt       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_belt : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_pitch_belt : int    NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_belt  : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ var_total_accel_belt : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ avg_roll_belt      : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_belt    : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_belt      : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_belt     : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_belt   : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_belt     : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_belt       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_belt     : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_belt       : num   NA NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_belt_x       : num    0 0.02 0 0.02 0.02 0.02 0.02 0.02 0.02 0.03 ...
## $ gyros_belt_y       : num    0 0 0 0 0.02 0 0 0 0 0 ...
## $ gyros_belt_z       : num   -0.02 -0.02 -0.02 -0.03 -0.02 -0.02 -0.02 -0.02 -0.02 0 ...
## $ accel_belt_x       : int   -21 -22 -20 -22 -21 -21 -22 -22 -20 -21 ...
## $ accel_belt_y       : int    4 4 5 3 2 4 3 4 2 4 ...
## $ accel_belt_z       : int   22 22 23 21 24 21 21 21 24 22 ...
## $ magnet_belt_x      : int   -3 -7 -2 -6 -6 0 -4 -2 1 -3 ...
## $ magnet_belt_y      : int   599 608 600 604 600 603 599 603 602 609 ...

```

```

## $ magnet_belt_z      : int  -313 -311 -305 -310 -302 -312 -311 -313 -312 -308 ...
## $ roll_arm           : num  -128 -128 -128 -128 -128 -128 -128 -128 -128 -128 ...
## $ pitch_arm          : num   22.5 22.5 22.5 22.1 22.1 22 21.9 21.8 21.7 21.6 ...
## $ yaw_arm            : num  -161 -161 -161 -161 -161 -161 -161 -161 -161 -161 ...
## $ total_accel_arm    : int    34 34 34 34 34 34 34 34 34 34 ...
## $ var_accel_arm      : num    NA NA NA NA NA NA NA NA NA NA ...
## $ avg_roll_arm       : num    NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_roll_arm    : num    NA NA NA NA NA NA NA NA NA NA ...
## $ var_roll_arm       : num    NA NA NA NA NA NA NA NA NA NA ...
## $ avg_pitch_arm      : num    NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_pitch_arm   : num    NA NA NA NA NA NA NA NA NA NA ...
## $ var_pitch_arm      : num    NA NA NA NA NA NA NA NA NA NA ...
## $ avg_yaw_arm        : num    NA NA NA NA NA NA NA NA NA NA ...
## $ stddev_yaw_arm     : num    NA NA NA NA NA NA NA NA NA NA ...
## $ var_yaw_arm        : num    NA NA NA NA NA NA NA NA NA NA ...
## $ gyros_arm_x        : num    0 0.02 0.02 0.02 0 0.02 0 0.02 0.02 0.02 ...
## $ gyros_arm_y        : num    0 -0.02 -0.02 -0.03 -0.03 -0.03 -0.03 -0.02 -0.03 -0.03 ...
## $ gyros_arm_z        : num   -0.02 -0.02 -0.02 0.02 0 0 0 0 -0.02 -0.02 ...
## $ accel_arm_x        : int   -288 -290 -289 -289 -289 -289 -289 -289 -288 -288 ...
## $ accel_arm_y        : int    109 110 110 111 111 111 111 111 109 110 ...
## $ accel_arm_z        : int   -123 -125 -126 -123 -123 -122 -125 -124 -122 -124 ...
## $ magnet_arm_x       : int   -368 -369 -368 -372 -374 -369 -373 -372 -369 -376 ...
## $ magnet_arm_y       : int    337 337 344 344 337 342 336 338 341 334 ...
## $ magnet_arm_z       : int    516 513 513 512 506 513 509 510 518 516 ...
## $ kurtosis_roll_arm  : num    NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_arm : num    NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_arm   : num    NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_roll_arm  : num    NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_pitch_arm : num    NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_arm   : num    NA NA NA NA NA NA NA NA NA NA ...
## $ max_roll_arm       : num    NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_arm      : num    NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_arm        : int     NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_arm       : num    NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_arm      : num    NA NA NA NA NA NA NA NA NA NA ...

```

```
## $ min_yaw_arm : int NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_arm : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_pitch_arm : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_yaw_arm : int NA NA NA NA NA NA NA NA NA NA NA ...
## $ roll_dumbbell : num 13.1 13.1 12.9 13.4 13.4 ...
## $ pitch_dumbbell : num -70.5 -70.6 -70.3 -70.4 -70.4 ...
## $ yaw_dumbbell : num -84.9 -84.7 -85.1 -84.9 -84.9 ...
## $ kurtosis_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ kurtosis_yaw_dumbbell : logi NA NA NA NA NA NA NA ...
## $ skewness_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ skewness_yaw_dumbbell : logi NA NA NA NA NA NA NA ...
## $ max_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ max_yaw_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_pitch_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ min_yaw_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## $ amplitude_roll_dumbbell : num NA NA NA NA NA NA NA NA NA NA NA ...
## [list output truncated]
```

As one can see in the “str(training)”, the dataset has a lot of NAs. Let’s first clean the data basically removing the NAs, the IDs and the zero variability columns.

```

# Starting the Cleaning Process
nzvCol <- nearZeroVar(training)
training <- training[,-nzvCol]

# Since we have lots of variables, remove any with NA's or have empty strings, and the one's that are not predictors variables
filterData <- function(idf) {
  idx.keep <- !sapply(idf, function(x) any(is.na(x)))
  idf <- idf[, idx.keep]
  idx.keep <- !sapply(idf, function(x) any(x==""))
  idf <- idf[, idx.keep]

  # Remove the columns that aren't the predictor variables
  col.rm <- c("X", "user_name", "raw_timestamp_part_1", "raw_timestamp_part_2",
             "cvtd_timestamp", "new_window", "num_window")
  idx.rm <- which(colnames(idf) %in% col.rm)
  idf <- idf[, -idx.rm]
  return(idf)
}

training <- filterData(training)
finalTrainingDS <- training
dim(finalTrainingDS)

```

```
## [1] 19622    53
```

```

# Now let's perform the same cleaning process to the testing dataset as well
nzvCol <- nearZeroVar(testing)
testing <- testing[,-nzvCol]
testing <- filterData(testing)
finalTestingDS <- testing
dim(finalTestingDS)

```

```
## [1] 20 53
```

Data Partitioning

Now we'll partition the data into training and testing datasets.

```
inTrain <- createDataPartition(y=finalTrainingDS$classe, p=0.85, list=FALSE)
newTraining <- finalTrainingDS[inTrain, ]
newTesting <- finalTrainingDS[-inTrain, ]
dim(newTraining); dim(newTesting)
```

```
## [1] 16680 53
```

```
## [1] 2942 53
```

Model Selection

We'll run some simulations with different machine learning algorithms. We'll use Random Forest, Decision Trees, Naive Bayes, Linear Discriminant Analysis and Generalized Boosted Regression Model. Besides this we'll check if using Principal Component Analysis also generates a good basis for prediction.

Note: from the referenced paper we know the AdaBoost algo yields something better than 99.5% accuracy. For this work we'll consider "good" anything higher than 98%.

Training Control Parameters

```
#Some fitting params - Repeated 5 Cross Validations
fitControl <- trainControl(method="repeatedcv", number=5, repeats=1, verboseIter=FALSE)
```

Predicting models with PCA pre-processing

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.1156
##	2	1.5376	nan	0.1000	0.0801
##	3	1.4873	nan	0.1000	0.0828
##	4	1.4358	nan	0.1000	0.0559
##	5	1.3994	nan	0.1000	0.0562
##	6	1.3642	nan	0.1000	0.0432
##	7	1.3357	nan	0.1000	0.0422
##	8	1.3091	nan	0.1000	0.0360
##	9	1.2857	nan	0.1000	0.0321
##	10	1.2651	nan	0.1000	0.0361
##	20	1.1047	nan	0.1000	0.0178
##	40	0.9307	nan	0.1000	0.0092
##	60	0.8290	nan	0.1000	0.0057
##	80	0.7507	nan	0.1000	0.0043
##	100	0.6865	nan	0.1000	0.0034
##	120	0.6333	nan	0.1000	0.0020
##	140	0.5887	nan	0.1000	0.0032
##	150	0.5666	nan	0.1000	0.0022

Predicting Models without PCA

##	Iter	TrainDeviance	ValidDeviance	StepSize	Improve
##	1	1.6094	nan	0.1000	0.2331
##	2	1.4621	nan	0.1000	0.1616
##	3	1.3604	nan	0.1000	0.1254
##	4	1.2810	nan	0.1000	0.1118
##	5	1.2118	nan	0.1000	0.0919
##	6	1.1536	nan	0.1000	0.0734
##	7	1.1071	nan	0.1000	0.0591
##	8	1.0689	nan	0.1000	0.0663
##	9	1.0283	nan	0.1000	0.0505
##	10	0.9964	nan	0.1000	0.0565
##	20	0.7627	nan	0.1000	0.0241
##	40	0.5382	nan	0.1000	0.0106
##	60	0.4136	nan	0.1000	0.0064
##	80	0.3320	nan	0.1000	0.0043
##	100	0.2707	nan	0.1000	0.0036
##	120	0.2275	nan	0.1000	0.0023
##	140	0.1923	nan	0.1000	0.0019
##	150	0.1781	nan	0.1000	0.0029

For the sake of comparison, we can see the overall indicators for each prediction model in the table below:

```
# Analyzing the results
resultingDataTable <- rbind(cm_tree$overall, cm_tree_pca$overall, cm_lda$overall, cm_lda_pca$overall,
                           cm_nb$overall, cm_nb_pca$overall, cm_rf$overall, cm_rf_pca$overall,
                           cm_gbm$overall, cm_gbm_pca$overall)
rownames(resultingDataTable) <- c("Tree", "Tree w/ PCA", "LDA", "LDA w/ PCA",
                                "Naive Baeyes", "Naive Bayes w/ PCA", "Random Forest", "Random Forest w/ PCA",
                                "GBM", "GBM w/ PCA")

resultingDataTable
```

##	Accuracy	Kappa	AccuracyLower	AccuracyUpper
## Tree	0.5033990	0.3516115	0.4851663	0.5216250
## Tree w/ PCA	0.3793338	0.1669767	0.3617577	0.3971502
## LDA	0.7046227	0.6257193	0.6877698	0.7210677
## LDA w/ PCA	0.5275323	0.4010752	0.5093025	0.5457073
## Naive Baeyes	0.7474507	0.6766211	0.7313398	0.7630681
## Naive Bayes w/ PCA	0.6451394	0.5507213	0.6275423	0.6624473
## Random Forest	0.9955812	0.9944106	0.9924556	0.9976452
## Random Forest w/ PCA	0.9816451	0.9767775	0.9761180	0.9861819
## GBM	0.9629504	0.9531219	0.9554790	0.9694816
## GBM w/ PCA	0.8314072	0.7864595	0.8173808	0.8447711

##	AccuracyNull	AccuracyPValue	McnemarPValue
## Tree	0.5105370	7.861069e-01	NaN
## Tree w/ PCA	0.7396329	1.000000e+00	NaN
## LDA	0.3031951	0.000000e+00	1.237497e-32
## LDA w/ PCA	0.3072740	1.260017e-135	3.489996e-38
## Naive Baeyes	0.3769545	0.000000e+00	4.565690e-49
## Naive Bayes w/ PCA	0.2970768	0.000000e+00	7.462410e-09
## Random Forest	0.2848402	0.000000e+00	NaN
## Random Forest w/ PCA	0.2865398	0.000000e+00	NaN
## GBM	0.2872196	0.000000e+00	1.097292e-02
## GBM w/ PCA	0.2984364	0.000000e+00	9.194940e-14

Out-Of-Sample Error

Our out-of-sample error can be found using the (1 - Testing Accuracy), as evaluated below (for Random Forest Algo).

```
oos_error <- 1 - cm_rf$overall[1]
print(paste("Out of Error Sample is: ", oos_error * 100, "%"))
```

```
## [1] "Out of Error Sample is: 0.441876274643105 %"
```

Conclusion

From the resulting table we can see the Random Forest algorithm yields a better result, and that using Principal Component Analysis, with a threshold of 99% of the variance would decrease of accuracy in about (aprox) 2% points. Besides this, comparing our results with the original paper results, we can see we have reach a very good prediction accuracy using Random Forest algorithm.

Appendix

Variable Importance

Just for sake of curiosity, lets check each variable importance, in the final model (random forest)

```
varImp(modelFitRF)
```

```
## rf variable importance
##
##    only 20 most important variables shown (out of 52)
##
##              Overall
## roll_belt      100.00
## yaw_belt       82.38
## magnet_dumbbell_z 70.10
## pitch_belt     66.28
## pitch_forearm  64.12
## magnet_dumbbell_y 61.37
## magnet_dumbbell_x 55.51
## roll_forearm   52.88
## accel_belt_z   47.39
## accel_dumbbell_y 44.42
## magnet_belt_z  43.88
## magnet_belt_y  43.05
## roll_dumbbell  40.15
## accel_dumbbell_z 36.47
## roll_arm       36.24
## accel_forearm_x 32.25
## accel_dumbbell_x 30.68
## yaw_dumbbell   30.04
## gyros_dumbbell_y 29.28
## magnet_forearm_z 28.86
```

Generating files to submit

Now we'll use the original testing dataset and the best model for predicting the values.

```
# This uses the code supplied by the class instructions
answers <- predict(modelFitRF, newdata=testing)
pml_write_files = function(x){
  n = length(x)
  for(i in 1:n){
    filename = paste0("problem_id_",i,".txt")
    write.table(x[i],file=filename,quote=FALSE,row.names=FALSE,col.names=FALSE)
  }
}
pml_write_files(answers)
```