

```
---
title: "Practical Machine Learning - Course Final Project"
author: "Alberto A. Caeiro Jr"
date: "February 16, 2015"
output: html_document
---
```

#### # Executive Summary & Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement – a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways.

The goal of this project is to predict the manner in which they did the exercise. This is the "classe" variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

#### # About the Data

The data for this project are available here:

. Training dataset :

"<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>"

. Testing dataset:

"<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>"

The data for this project come from this source:

<http://groupware.les.inf.puc-rio.br/har>. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment. So let's have a look on the dataset and on the classe variable.

#### # Data Processing

## Obtaining and Cleaning the data

```
```{r, message=FALSE}
```

# Enabling Multi Core for modeling processing

```
library(doMC)
```

```
  registerDoMC(cores = 2)
```

#Loading used libraries

```
library(caret);library(klaR); library(rpart)
```

```
library(randomForest); library(gbm)
```

#setting the seed for reproducible computation

```
set.seed(12345)
```

#setting the working directory folder

```
setwd("~/Developer/Data Science Specialization/Practical Machine
```

```
Learning/Project")
```

```
# loading both testing and training dataset (considering both files were  
already downloaded)
```

```
trainFile <- "./pml-training.csv"
```

```
training <- read.csv(file=trainFile, header=TRUE, sep="," ,
```

```
na.strings=c("NA", "#DIV/0!", ""))
```

```
testFile <- "./pml-testing.csv"
```

```
testing <- read.csv(file=testFile, header=TRUE, sep="," ,
```

```
na.strings=c("NA", "#DIV/0!", ""))
```

```
# Summary for the training predictors and outcome
```

```
str(training)
```

```
```
```

As one can see in the "str(training)", the dataset has a lot of NAs.  
Let's first clean the data basically removing the NAs, the IDs and the  
zero variability columns.

```
```{r}
```

```
# Starting the Cleaning Process
```

```
nzvCol <- nearZeroVar(training)
```

```
training <- training[, -nzvCol]
```

```
# Since we have lots of variables, remove any with NA's or have empty  
strings, and the one's that are not predictors variables
```

```
filterData <- function(idf) {
```

```
  idx.keep <- !sapply(idf, function(x) any(is.na(x)))
```

```
  idf <- idf[, idx.keep]
```

```
  idx.keep <- !sapply(idf, function(x) any(x==""))
```

```
  idf <- idf[, idx.keep]
```

```
  # Remove the columns that aren't the predictor variables
```

```
  col.rm <- c("X", "user_name", "raw_timestamp_part_1",
```

```
"raw_timestamp_part_2",
```

```
            "cvtd_timestamp", "new_window", "num_window")
```

```
  idx.rm <- which(colnames(idf) %in% col.rm)
```

```
  idf <- idf[, -idx.rm]
```

```
  return(idf)
```

```
}
```

```
training <- filterData(training)
```

```
finalTrainingDS <- training
```

```
dim(finalTrainingDS)
```

```
# Now let's perform the same cleaning process to the testing dataset as  
well
```

```
nzvCol <- nearZeroVar(testing)
```

```
testing <- testing[, -nzvCol]
```

```
testing <- filterData(testing)
```

```
finalTestingDS <- testing
```

```
dim(finalTestingDS)
```

```
```
```

```
## Data Partitioning
```

```
Now we'll partition the data into training and testing datasets.
```

```
```{r}
```

```
inTrain <- createDataPartition(y=finalTrainingDS$classe, p=0.85,
```

```
list=FALSE)
```

```

newTraining <- finalTrainingDS[inTrain, ]
newTesting <- finalTrainingDS[-inTrain, ]
dim(newTraining); dim(newTesting)
```

## Model Selection
We'll run some simulations with different machine learning algorithms.
We'll use Random Forest, Decision Trees, Naive Bayes, Linear Discriminant
Analysis and Generalized Boosted Regression Model. Besides this we'll
check if using Principal Component Analysis also generates a good basis
for prediction.

. Note: from the referenced paper we know the AdaBoost algo yields 99,9%
accuracy. For this work we'll consider "good" anything higher than 98%.

## Training Control Parameters
```{r}
#Some fitting params - Repeated 5 Cross Validations
fitControl <- trainControl(method="repeatedcv", number=5, repeats=1,
verboseIter=FALSE)
```

## Predicting models with PCA pre-processing
```{r, cache=TRUE, warning=FALSE, message=FALSE, echo=FALSE }
prePro <- preProcess(newTraining[, -53], method="pca", tresh=0.99)
newTrainingPCA <- predict(prePro, newTraining[, -53]) ; newTestingPCA <-
predict(prePro, newTesting[, -53])
```

# Decision Trees
modelFitTreePCA <- train(newTraining$classe ~ . , method="rpart",
trControl=fitControl, data=newTrainingPCA)
cm_tree_pca <- confusionMatrix(newTesting$classe,
predict(modelFitTreePCA, newdata=newTestingPCA))

# Linear Discriminant Analysis
modelFitLDAPCA <- train(newTraining$classe ~ . , method="lda",
trControl=fitControl, data=newTrainingPCA)
cm_lda_pca <- confusionMatrix(newTesting$classe, predict(modelFitLDAPCA,
newdata=newTestingPCA))

# Naive Bayes
modelFitNBPCA <- train(newTraining$classe ~ . , method="nb",
trControl=fitControl, data=newTrainingPCA)
cm_nb_pca <- confusionMatrix(newTesting$classe, predict(modelFitNBPCA,
newdata=newTestingPCA))

# Random Forest
modelFitRFPCA <- train(newTraining$classe ~ . , method="rf",
data=newTrainingPCA)
cm_rf_pca <- confusionMatrix(newTesting$classe, predict(modelFitRFPCA,
newdata=newTestingPCA))

# Generalized Boosted Regression Modeling
modelFitGBMPCA <- train(newTraining$classe ~ . , method="gbm",
trControl=fitControl, data=newTrainingPCA)
cm_gbm_pca <- confusionMatrix(newTesting$classe, predict(modelFitGBMPCA,
newdata=newTestingPCA))

```

```
```
```

```
## Predicting Models without PCA
```{r, cache=TRUE, warning=FALSE, message=FALSE, echo=FALSE }
# Decision Trees
modelFitTree <- train(classe ~ . , method="rpart", data=newTraining,
trControl=fitControl)
cm_tree <- confusionMatrix(newTesting$classe, predict(modelFitTree,
newdata=newTesting))

# Linear Discriminant Analysis
modelFitLDA <- train(classe ~ ., method="lda", data=newTraining,
trControl=fitControl)
cm_lda <- confusionMatrix(newTesting$classe, predict(modelFitLDA,
newdata=newTesting))

#Naive Bayes
modelFitNB <- train(classe ~ ., method="nb", data=newTraining,
trControl=fitControl)
cm_nb <- confusionMatrix(newTesting$classe, predict(modelFitNB,
newdata=newTesting))

#Random Forest
modelFitRF <- train(classe ~ . , method="rf", data=newTraining)
cm_rf <- confusionMatrix(newTesting$classe, predict(modelFitRF,
newdata=newTesting))

#Generalized Boosted Regression Modeling
modelFitGBM <- train(classe ~ ., method="gbm", data=newTraining,
trControl=fitControl)
cm_gbm <- confusionMatrix(newTesting$classe, predict(modelFitGBM,
newdata=newTesting))
```
```

For the sake of comparison, we can see the overall indicators for each prediction model in the table below:

```
```{r}
# Analyzing the results
resultingDataTable <- rbind(cm_tree$overall, cm_tree_pca$overall,
cm_lda$overall, cm_lda_pca$overall,
cm_nb$overall, cm_nb_pca$overall, cm_rf$overall,
cm_rf_pca$overall,
cm_gbm$overall, cm_gbm_pca$overall)
rownames(resultingDataTable) <- c("Tree", "Tree w/ PCA", "LDA", "LDA w/
PCA",
"Naive Baeyes", "Naive Bayes w/ PCA",
"Random Forest", "Random Forest w/ PCA",
"GBM", "GBM w/ PCA")
```

```
resultingDataTable
```

```
```
```

```
## Out-Of-Sample Error
```

Our out-of-sample error can be found using the (1 - Testing Accuracy), as evaluated below (for Random Forest Algo).

```
```{r}
oos_error <- 1 - cm_rf$overall[1]
```

```
print(paste("Out of Error Sample is: ", oos_error * 100, "%"))
```
```

#### # Conclusion

From the resulting table we can see the Random Forest algorithm yields a better result, and that using Principal Component Analysis, with a threshold of 99% of the variance would decrease of accuracy in about (aprox) 2% points.

Besides this, comparing our results with the original paper results, we can see we have reach a very good prediction accuracy using Random Forest algorithm.

#### # Appendix

##### ## Variable Importance

Just for sake of curiosity, lets check each variable importance, in the final model (random forest)

```
```{r}
varImp(modelFitRF)
```
```

#### # Generating files to submitt

Now we'll use the original testing dataset and the best model for predicting the values.

```
```{r}
# This uses the code supplied by the class instructions
answers <- predict(modelFitRF, newdata=testing)
write_files = function(x){
  n = length(x)
  for(i in 1:n){
    filename = paste0("problem_id_",i,".txt")

write.table(x[i],file=filename,quote=FALSE,row.names=FALSE,col.names=FALSE)

  }
}
write_files(answers)
```
```